

# ROSMOD: A Domain Specific Tool-suite (DSTS) for Distributed CPS

William Emfinger, Pranav Srinivas Kumar, and Gabor Karsai

Kumar, P.S.; Emfinger, W.; Karsai, G.; Watkins, D.; Gasser, B.; Anilkumar, A. ROSMOD: A Toolsuite for Modeling, Generating, Deploying, and Managing Distributed Real-time Component-based Software using ROS. *Electronics* **2016**, *5*, 53.

This work was supported by DARPA under contract NNA11AB14C and USAF/AFRL under Cooperative Agreement FA8750-13-2-0050, and by the National Science Foundation (CNS-1035655). Any opinions, findings, and conclusions or recommendations expressed in this material are those of the author(s) and do not necessarily reflect the views of DARPA, USAF/AFRL, or the NSF.

## Outline

- Motivation
- Solution: ROSMOD
  - Why WebGME
- Design Principles
  - Meta-model / language
  - Functionality
  - Interface
- Implementation / Demo
  - How to structure the Meta-model
  - How to structure the WebGME components
- Lessons Learned



# Motivation

- Distributed CPS are hard to design, develop, analyze, deploy, and manage
- Integration of these key requirements into an IDE would make these processes
  - Easier to teach
  - Less error prone
  - Faster / More repeatable
- Some IDEs are heavy and complicated to install / set-up
  - Require training as well
  - Need maintaining to ensure proper versions and roll-out of updates/bug-fixes
  - Need to be cross-platform
- Not every system (or type of system) can or should conform to the same meta
  - Even within the same class of “Distributed CPS”
  - Not everyone wants or needs the context of very explicit / fine-grained network specification

# Solution

- Need an extensible, modular tool-suite that allows users / developers to create or swap feature sets for different deployments
- Need to be able to easily change the language associated with the modeling to tightly fit the class of systems being developed
- Need to make installation, training, and updating of the IDE easy for end-users
  - Especially when the IDE interacts with a lot of back-end infrastructure, e.g. for CI builds/tests or deployment
- All of these concerns leads us to:
  - WebGME: yes it can be better than eclipse.

# Design Principles: Meta (1/2)

- Need to capture only the abstract concepts related to the design and development of the class of systems in question
- Must balance between creating abstractions for everything (generalization to more systems) and implementation details (coupling to a technology / system)
  - Generalization is good, but can come at the cost of usability / user-friendliness
  - Implementation details can require proliferation of meta-models to account for variations in the implementation, but can improve usability / user-friendliness
- More than just **what** is captured in the meta is important; it matters **how** and **where** in the meta the concepts exist
  - This translates to *containment* / *inheritance* of the objects and how their *attributes* are represented
- **The design of the meta is the most important part.**

# Design Principles: Meta (2/2)

- When designing the meta you must also think about:
  - Is the meta something that should be extensible by users?
  - Should users even see the meta?
- The meta is the main user-interface a modeler will see, and the one they will interact with the most.
  - Good meta → happy users
  - Can be worked around with some good visualizers, but use sparingly
    - A visualizer / decorator only changes one aspect of the interaction, you'd need to change the other webgme components as well, e.g. *TreeBrowser*, *PartBrowser*, *AttributePanel*, etc.
- Finally, anything required for the class of systems that is not in the meta will not be available to the users but can be controlled in you components
  - E.g. the connection mechanisms to target nodes, or the middleware library in use for communication
  - Need to draw the line between the platform (infrastructure) and the model

# Design Principles: Functionality (1/2)

- Once you've designed the meta, you need to make sure its design facilitates the creation of models
  - Especially important for **sets** and **pointers**:
    - Do you want a pointer to be a pointer or a *connection*?
    - Is a set actually a set, or just a collection of connections?
  - Similarly for containment:
    - Just because the domain you're modeling has a hierarchy / tree-structure, doesn't mean your meta/model should. Heavily dependent on what kind of user-interaction you foresee.
- Now that you've designed your meta and made sure it's painless to make models, what else do you want your domain specific tool-suite (*DSTS*) to do?

# Design Principles: Functionality (2/2)

- Generally, you're wanting to
  - Get some parts of the model
  - Perform some transformation
  - (Rendering) -- visual, textual, etc.
  - (Update) -- write-back
- Do you want it as a **Plugin, Visualizer, Addon**? Depends on:
  - How you want the users to trigger and interact with it
  - What functional dependencies your functionality has (e.g. file-system, user-interaction, remote computing, etc.)
  - How many users/instances can be active?
- Also influenced by how reusable your code-base is and how dependent it is on your meta
  - Similar trade-off as meta design: modular, meta-agnostic code (generalization) versus meta-specific code that only works with your current meta model
  - Generalization is good (for everyone) if done right, but can be a challenge to achieve
  - Implementation specific code changes every time your meta changes (which may be frequently)

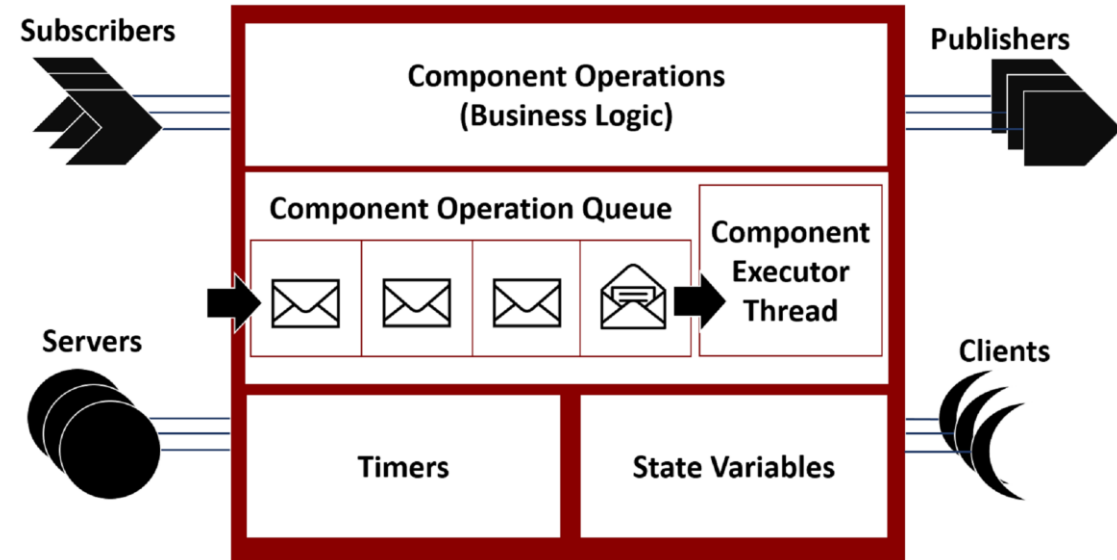


# Design Principles: Interface

- Final part (that you should do last)
  - But unfortunately you should think about it during the rest of the processes
- Visual styling of the meta/models
  - Icons
  - Decorators
  - Visualizers
  - Themes (color schemes)
- These improve quality of life, but are likely to change during development of the rest of the DSTS
  - As you're developing model transformation / analysis code, you're still likely to update the meta as you find elements you've forgotten or not represented in a manageable way

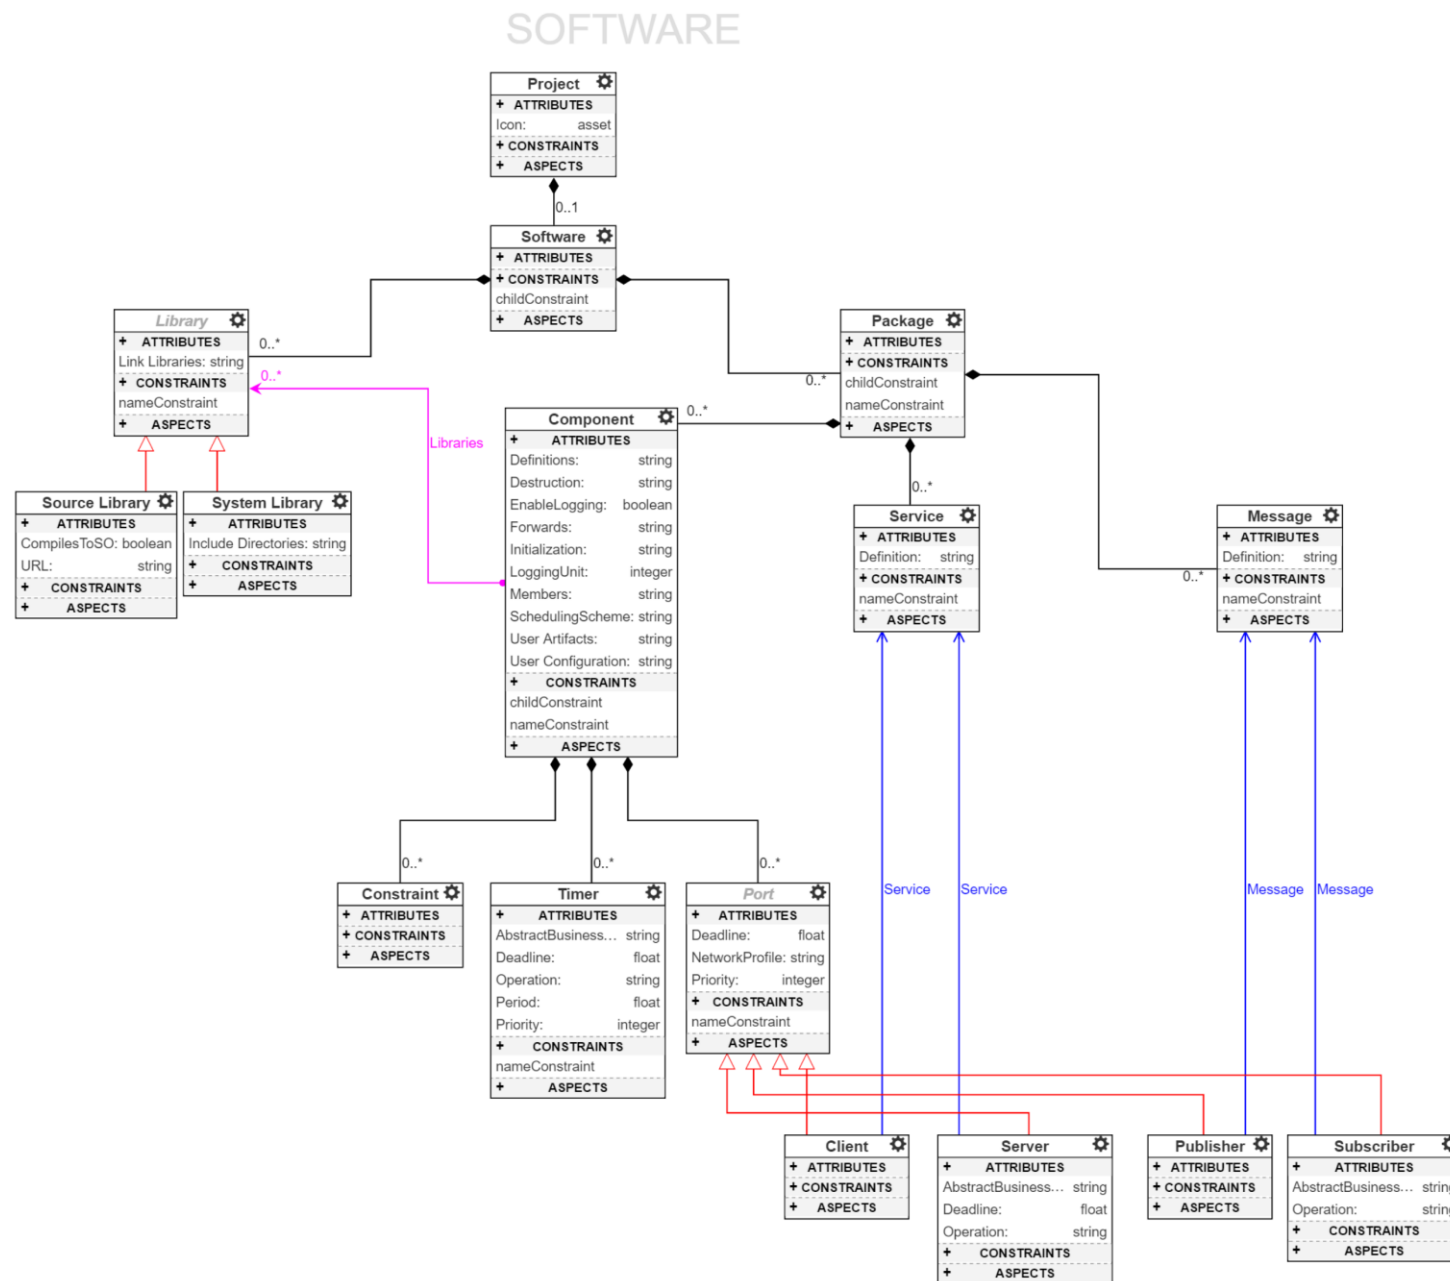
# ROSMOD Specifics

- Meta contains everything needed to specify
  - Distributed, component-based software
  - Networked embedded systems
  - Deployments of software onto hardware
  - Experiment executions and their results
- Plugins enable
  - Code generation/compilation
  - Functional (timing) model analysis
  - Documentation generation
  - Experiment deployment / execution
  - Experiment Teardown and results aggregation
- Visualizers enable
  - Project browser with relevant descriptions and identification
  - Deployment visual inspection and call chain tracing
  - Execution trace log visualization



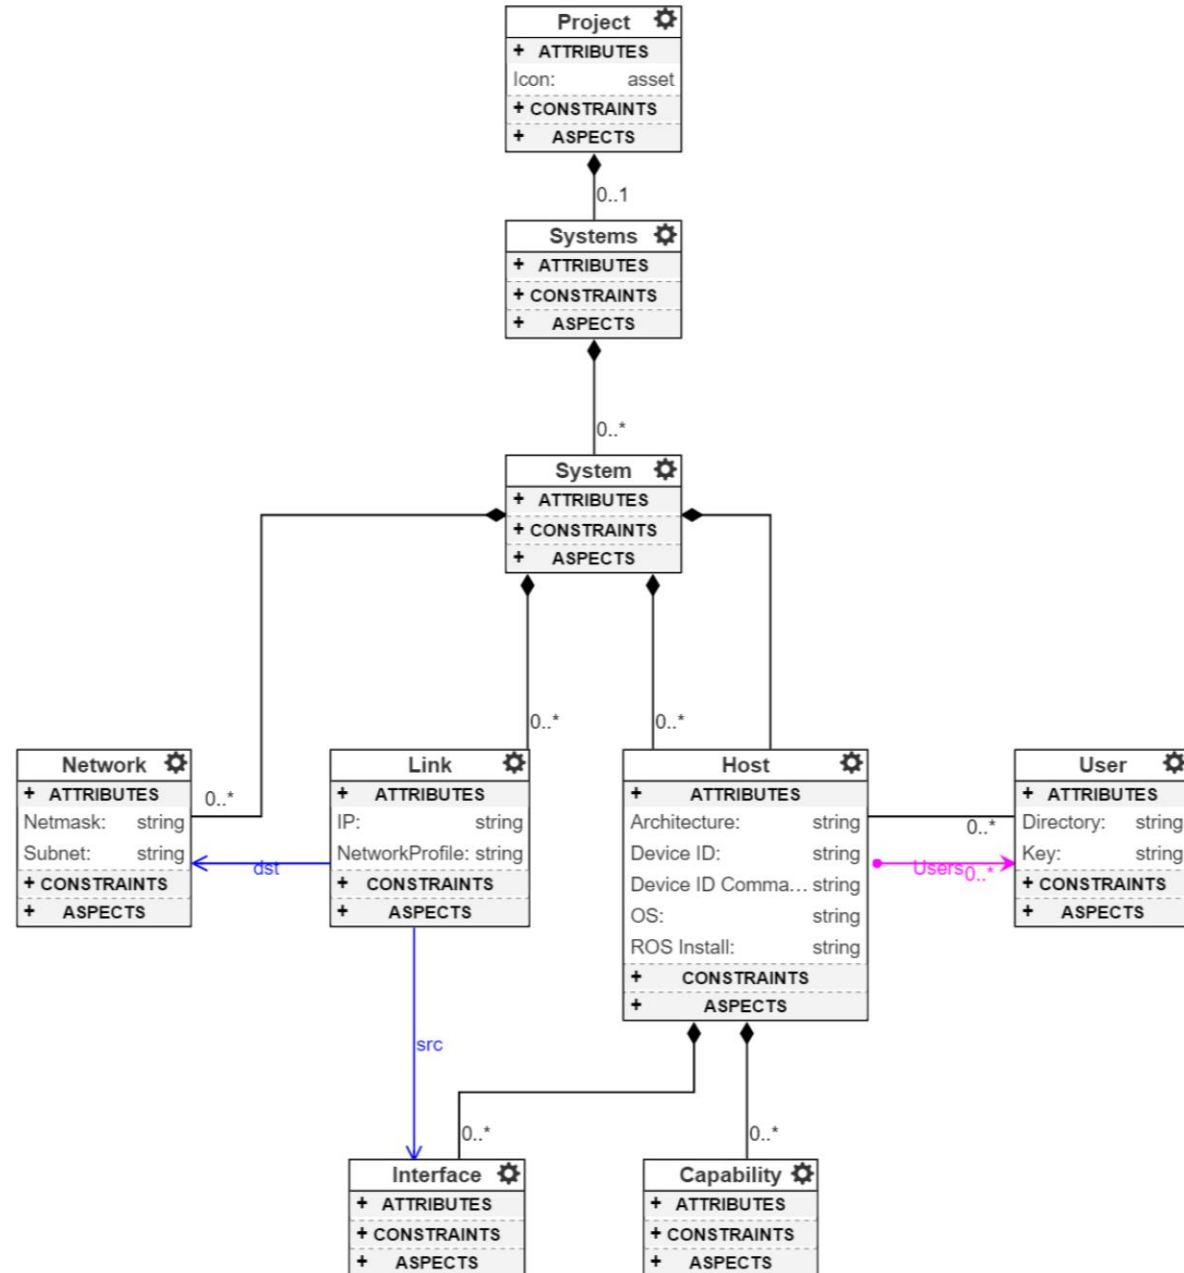
# Implementation: Meta

- Software contains generic concepts like libraries, operations, definitions, etc.
  - Mostly language agnostic; can be C/C++, python, or any other language which supports these concepts
- *Message* and *Service* pointers are not connection objects, don't drag / draw between objects to establish connection
  - Trade-off for ease of specification with simplicity and out-of-tree specification
  - Currently cross-cuts are not able to do what we need them to (and further complicate interface)



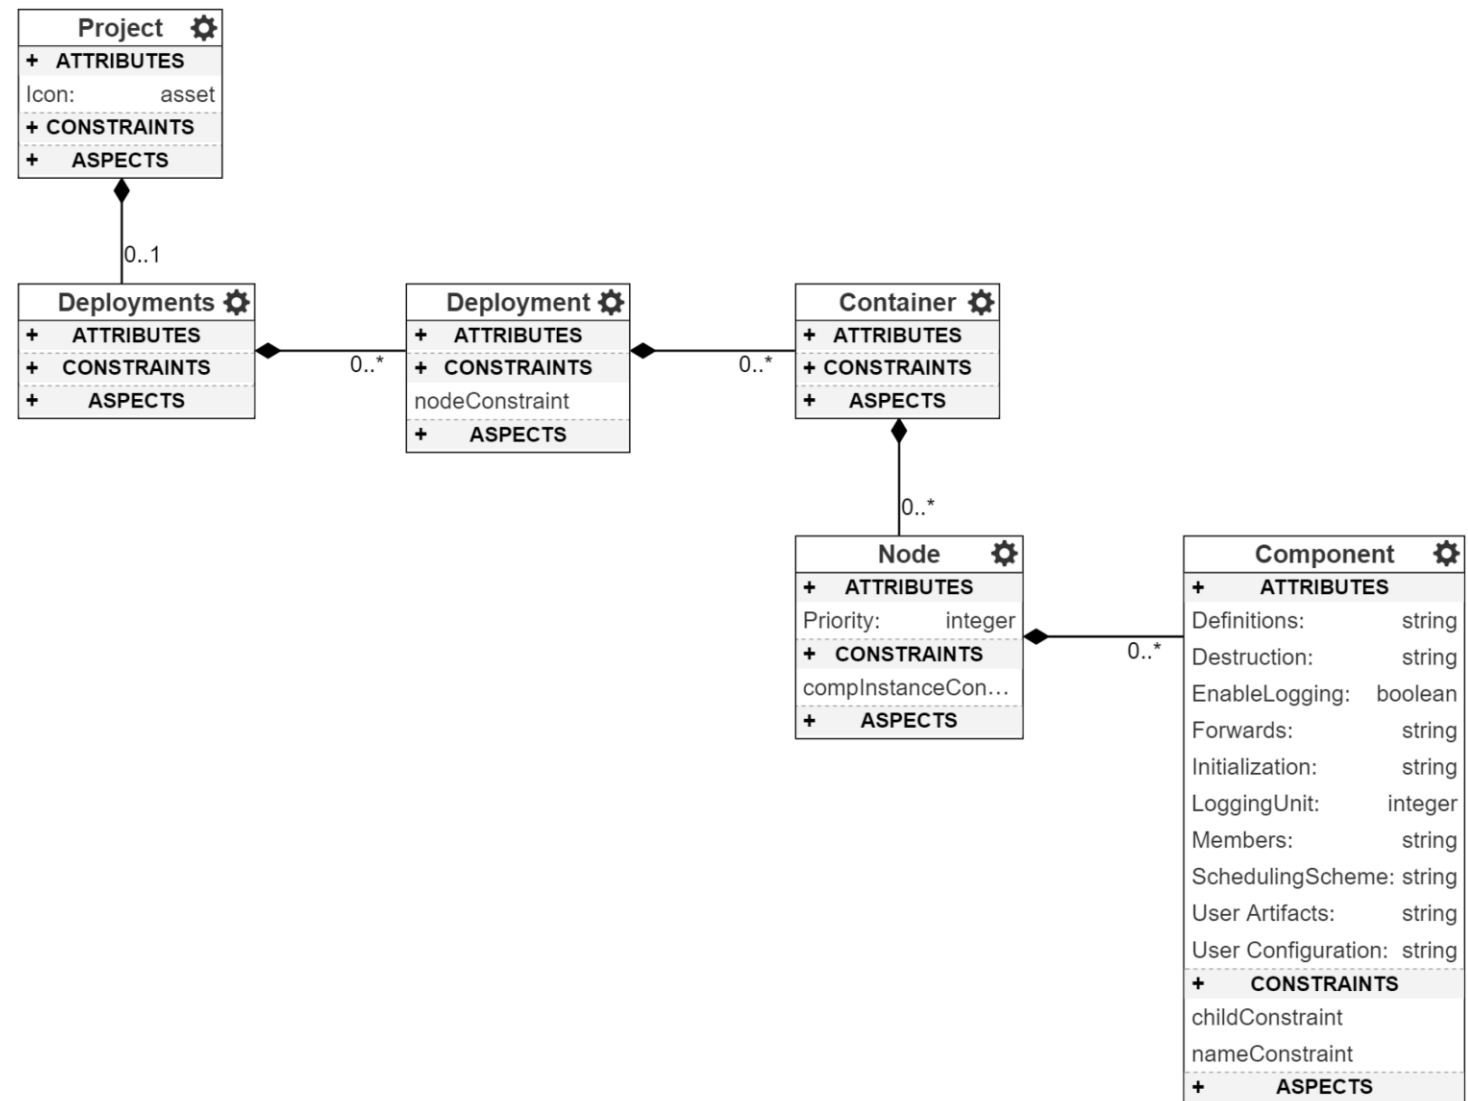
# Implementation: Meta

- This is a case of meta trying too hard to be like the domain it represents
  - Why did I have a *Link* between an *Interface* and a *Network*? I like Zelda, but no, I was just following the domain.
  - Means I have to create an *Interface* object inside a *Host* before drawing a link between its interface *port* and the network
    - Would be much simpler to just have the link between the *Host* and the *Network*
- Some attributes like *Key* are a little too implementation specific, but unavoidable



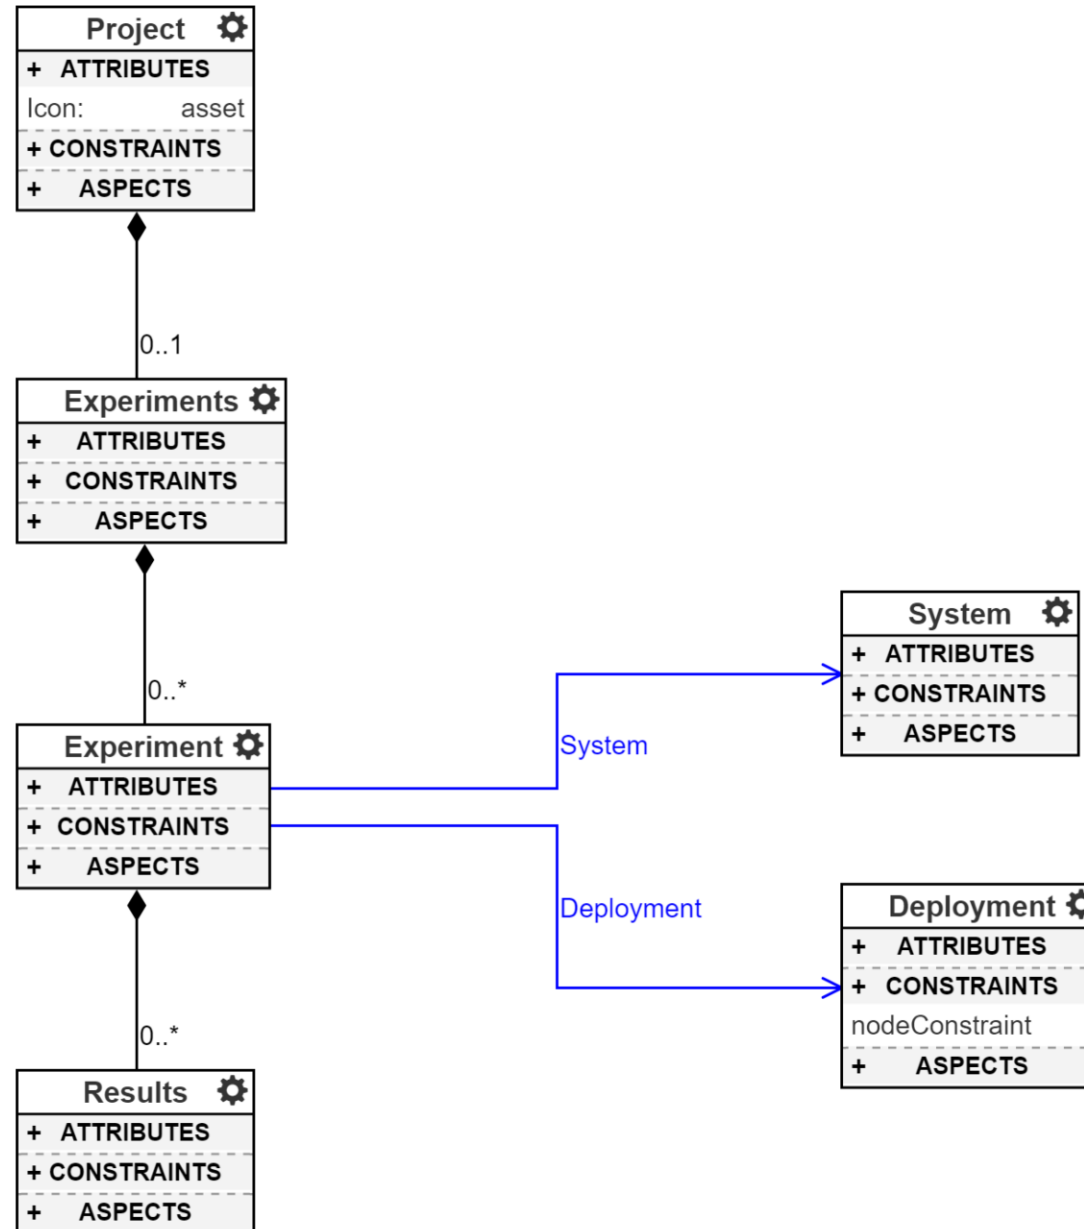
# Implementation: Meta

- Meant to be a generic collection of components in to nodes, nodes into containers
- Nice from a meta-modeling perspective, and lets us re-use these deployment specifications across systems
- But has been a source of confusion when training new users
  - Most don't really see the point of it
- Moral to this story: if you're going to generalize something, make sure it **simplifies** the concepts and the training required as much as possible
  - If you can't describe it simply, it may not be worth it.



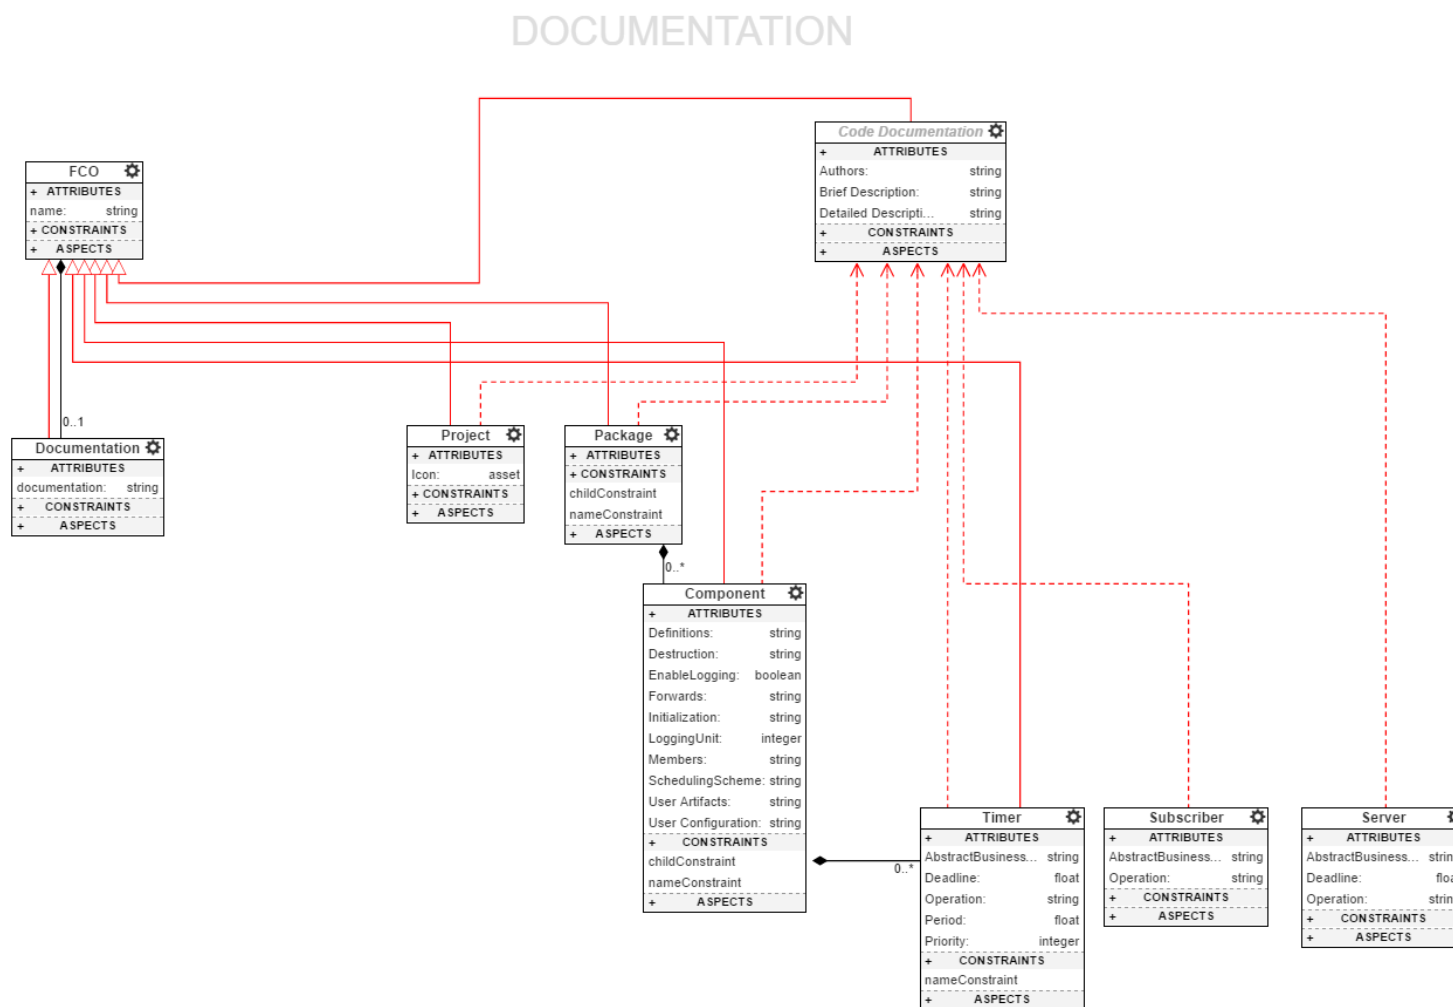
# Implementation: Meta

- Exhibits another downside of default pointer visualization, specification:
  - Can't easily see exactly which system or deployment an experiment is configured with.
- Need a way to visualize the composition of these two objects into one
  - Sounds like a job for a visualizer!
- However, the purpose of this relation was to have our infrastructure *automatically* calculate the mapping based on parameters that are **not modeled** because they are variable
  - E.g. current resource utilization or network connectivity



# Implementation: Meta

- Don't forget about **Mixins**, they allow multiple inheritance and nice object-oriented design of your meta!
- I can easily specify and change which objects have the same attributes and what the default attributes are
  - This also lets me generalize my plugins / visualizers a little more because now I can just tell other developers if they want to use my components for these purposes, they simply need to add a mixin with these attributes to their meta
- Only downside is that you can't see directly in the objects' attributes their inherited attributes



# Implementation: WebGME Components

- Project Browser
  - Allows users to easily navigate and select which project they want to work on
  - Gives more contextual information than standard decorators/visualizers
- Implemented without decorator
  - All in the visualizer
- Is this the right way to do this?
  - Depends on what your goal is 😊
- Getting some measure of *default behavior* is still a little tricky / undocumented
  - E.g. drag and drop

The screenshot displays the ROSMOD / Samples web application. The browser address bar shows the URL: <https://rosmod.rcps.isis.vanderbilt.edu/?project=ROSMOD%2BSamples&branch=master&node=%2Fv&visualizer=RootViz&layout=NewDefaultLayout>. The page features a navigation bar with 'ROSMOD' and 'ROSMOD / Samples' tabs, and a 'master' branch indicator. The main content area is divided into four panels, each representing a different project:

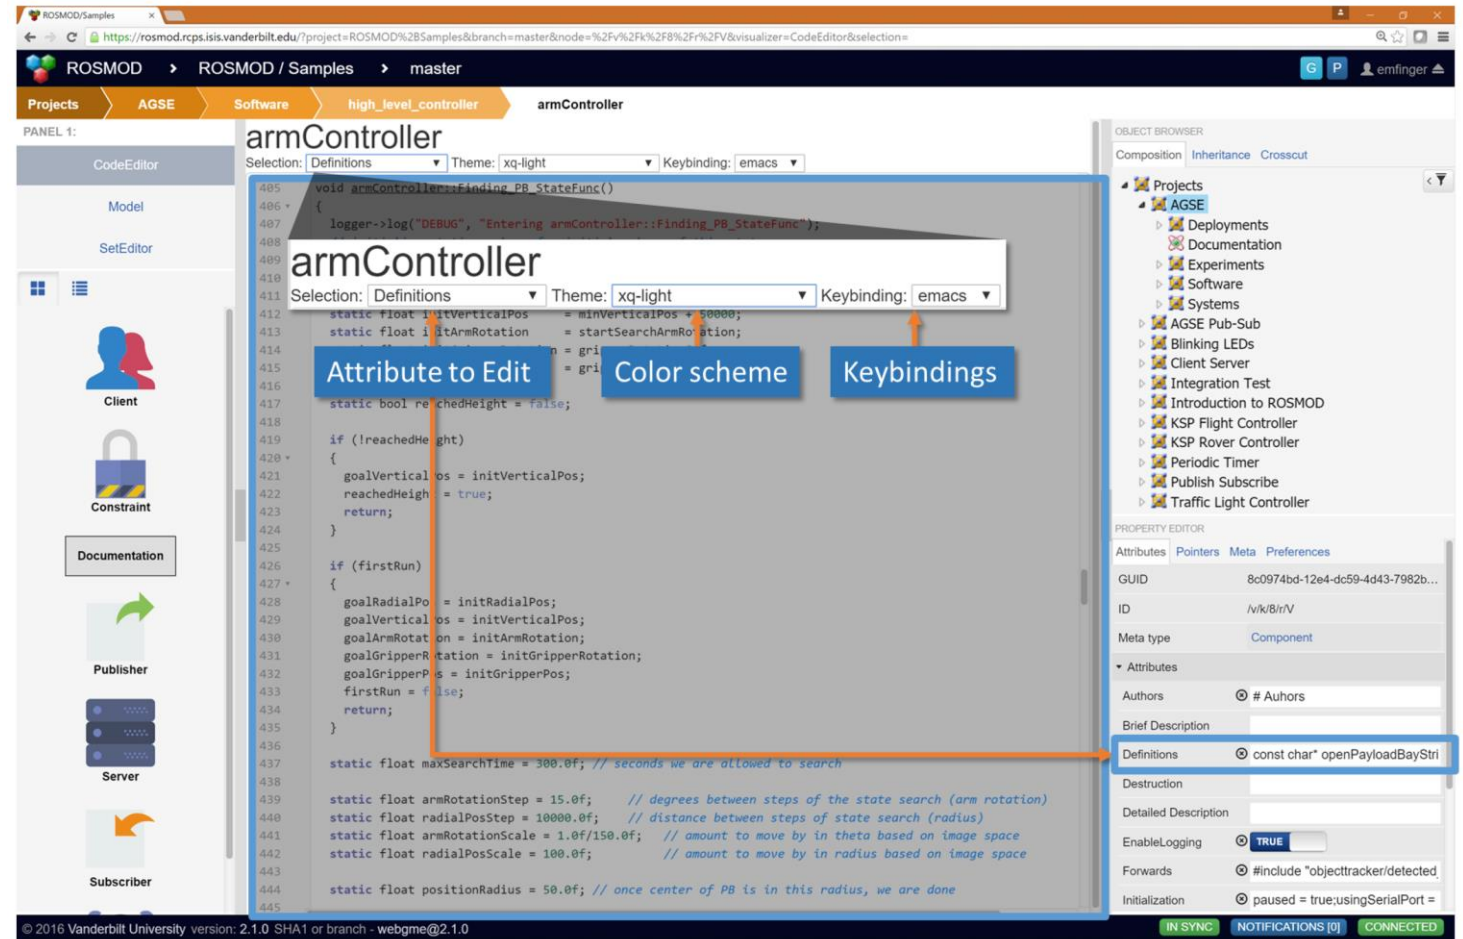
- Traffic Light Controller**: Includes an icon of a traffic light and a detailed description of its functionality and authors (William Emfinger).
- KSP Flight Controller**: Includes an icon of a rocket and a detailed description of its functionality and authors (Pranav Srinivas Kumar, William Emfinger).
- KSP Rover Controller**: Includes an icon of a rover and a detailed description of its functionality and authors (Pranav Srinivas Kumar, William Emfinger).
- AGSE**: Includes an icon of a robotic arm and a detailed description of its functionality and authors (Pranav Srinivas Kumar, William Emfinger).

On the left side, there is a 'Projects' panel with a 'RootViz' tab and a 'Documentation' icon. On the right side, there is an 'OBJECT BROWSER' panel showing a tree view of the project structure, including 'AGSE', 'Deployments', 'Documentation', 'Experiments', 'Software', 'Systems', 'Introduction to ROSMOD', 'KSP Flight Controller', 'KSP Rover Controller', and 'Traffic Light Controller'. Below the object browser is a 'PROPERTY EDITOR' panel with tabs for 'Attributes', 'Pointers', 'Meta', and 'Preferences'. The bottom of the page features a footer with copyright information, version details (2.5.1), and status indicators for 'IN SYNC', 'NOTIFICATIONS [0]', and 'CONNECTED'.



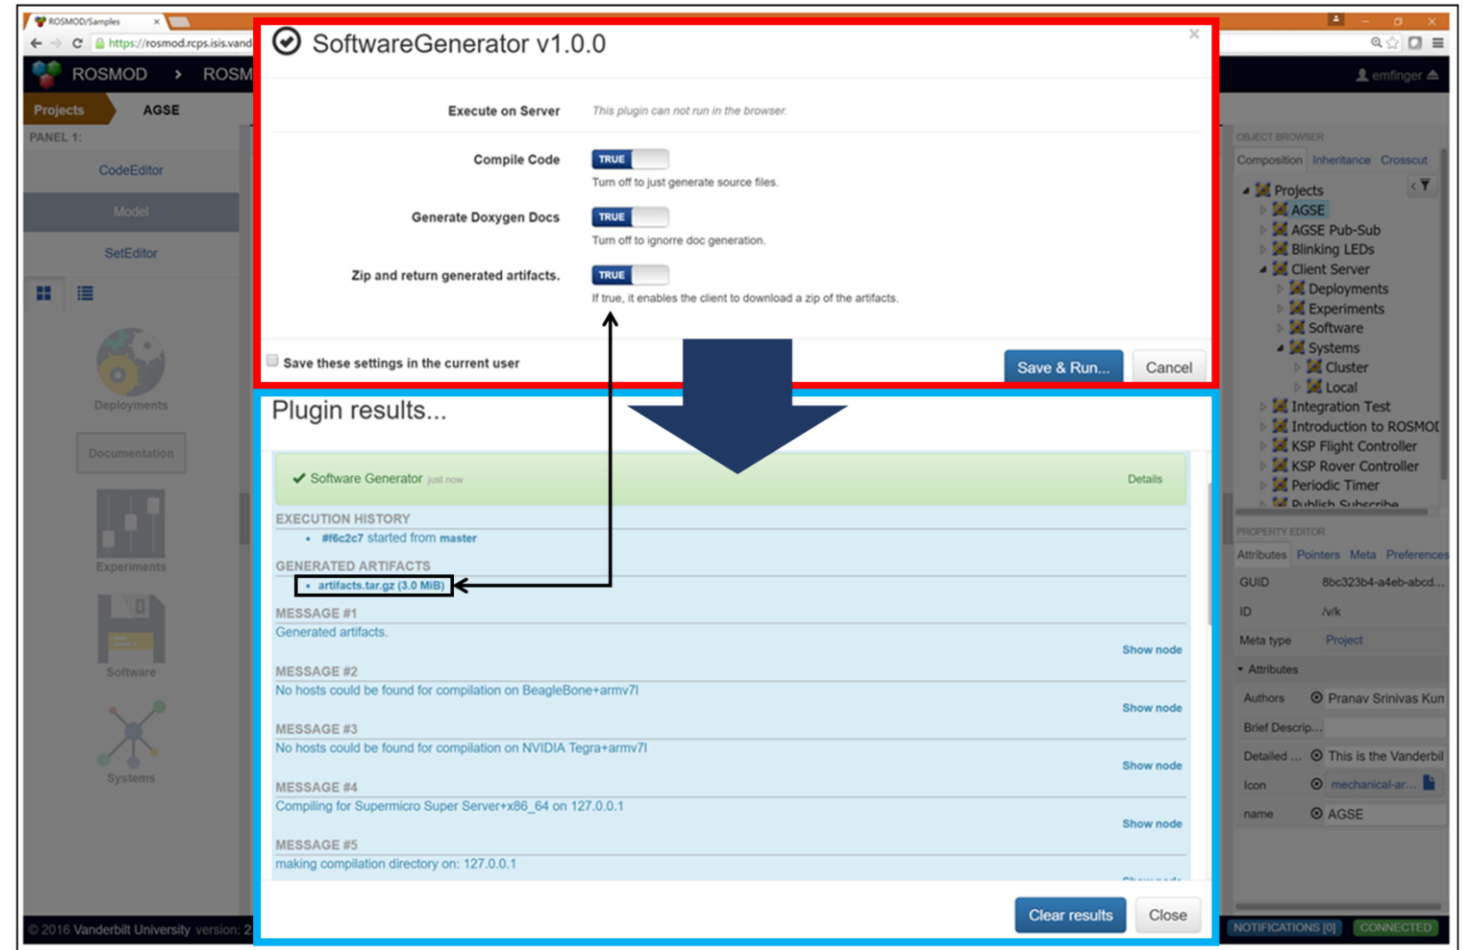
# Implementation: WebGME Components

- Code Editing in the browser; need features of an IDE:
  - Syntax highlighting
  - Code completion
  - Code folding
  - Theming ☺
  - Keybindings (don't start a vi(m) vs. emacs flame-war here...)
  - Multiple buffers
- The CodeEditor component is generic, and can be added to any WebGME project with configuration though `./config/components.json`



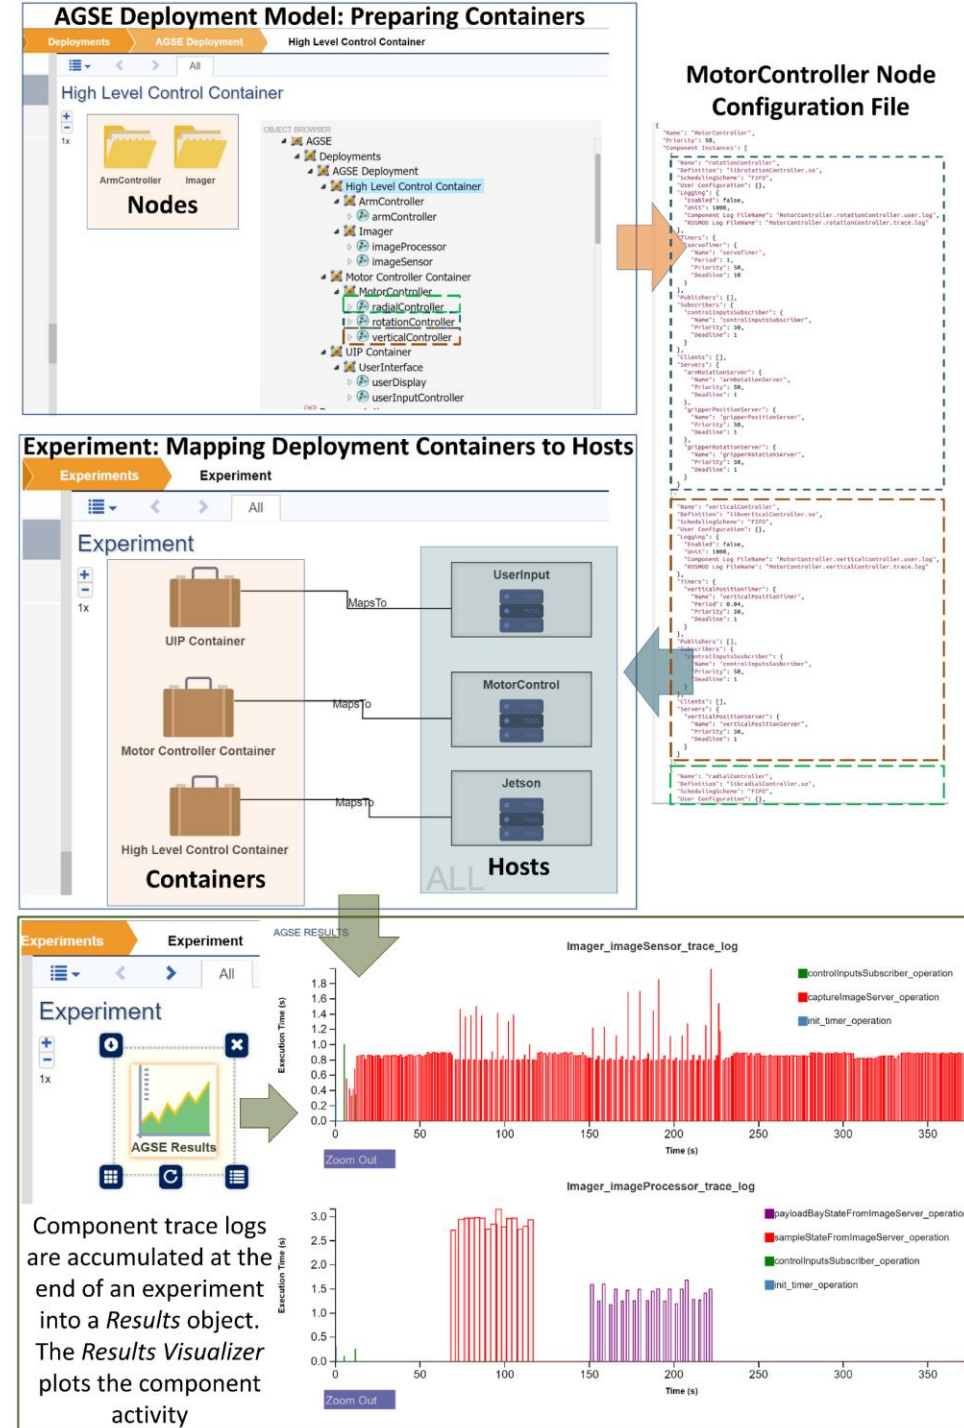
# Implementation: WebGME Components

- Software Generation and Compilation
  - All the code is either
    - Contained within the model,
    - Generatable from the model, or
    - Located in a repository as part of a library which is referenced in the model
  - This means that users don't have to touch the generated code → large quality of life improvement over previous systems we've used in the past
  - **Code templates should be as agnostic of the generation code as possible**
- Compilation (when required) must use the filesystem and call compilers which cannot run in the browser ☹
  - The compilation runs on the server, which is good for two reasons:
    - Users don't have to install or manage the compilers or their dependencies
    - Updates to the compilers can be managed in a centralized fashion by sys-admins



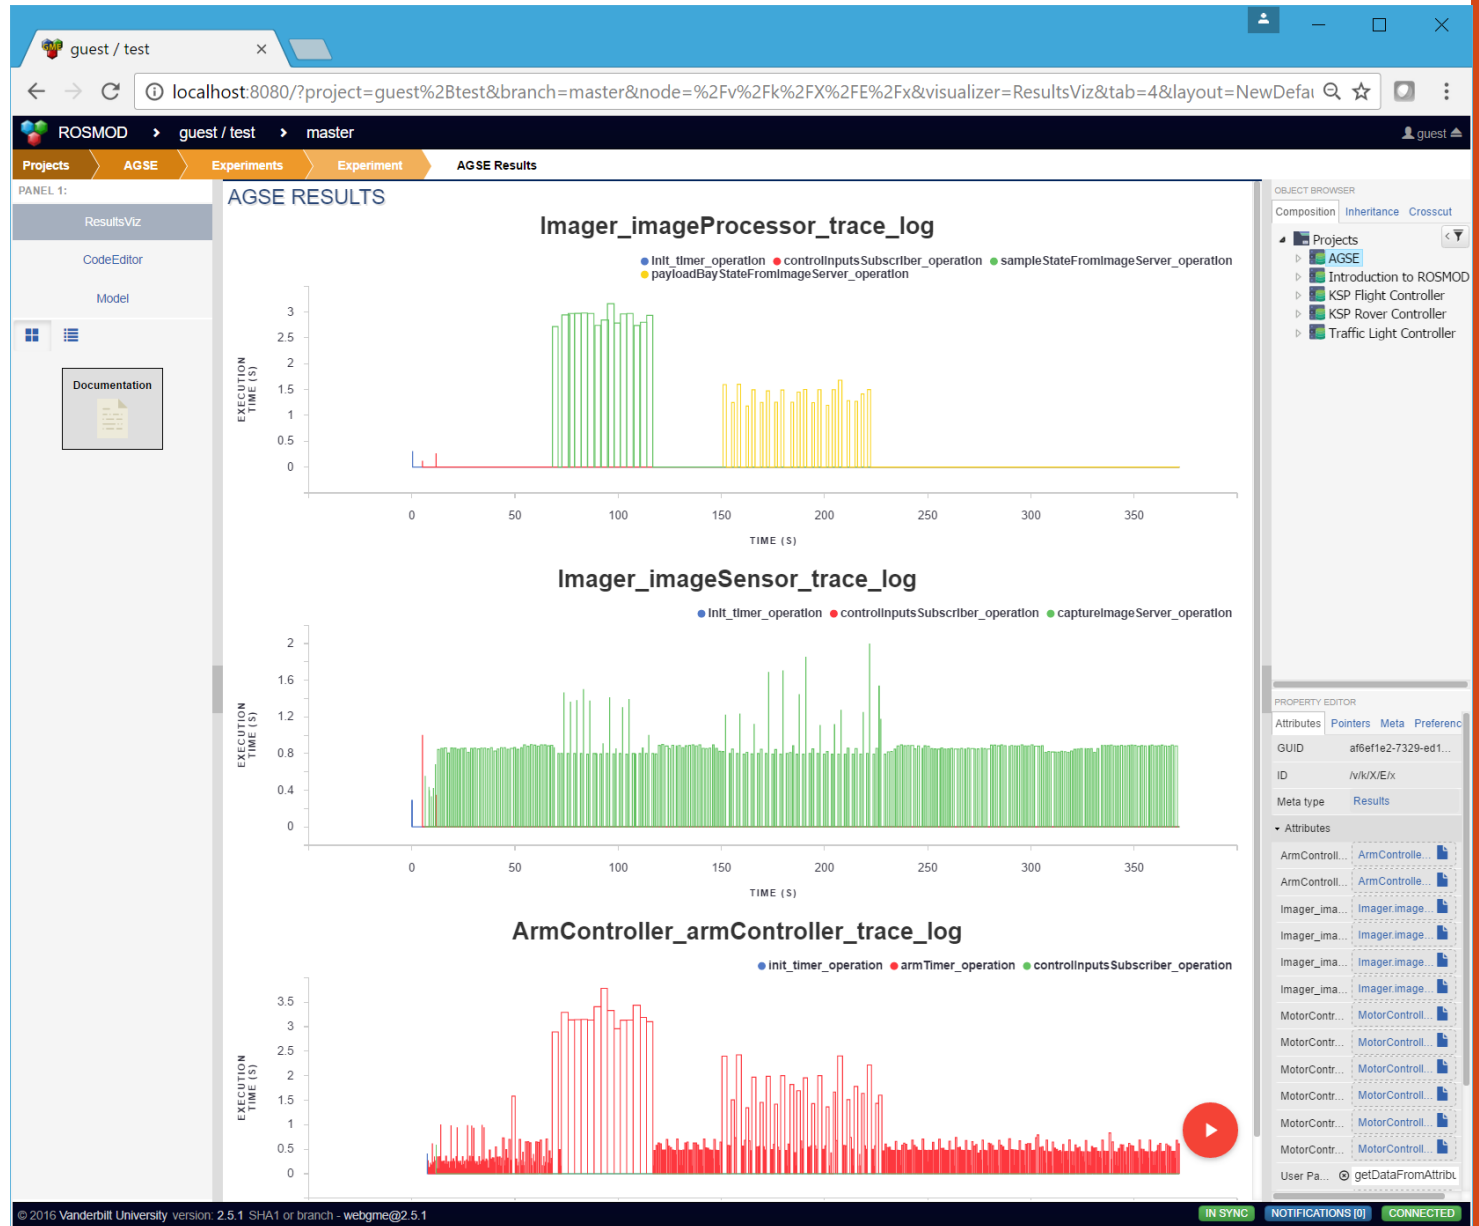
# Implementation: WebGME Components

- Experiment Deployment and Execution
  - Like the software generation/compilation, must run on the server, since it actually moves the binaries and configuration files over to the distributed systems
  - Automatically queries the systems described in the model to determine which have available resources for running the experiment
  - Updates the model to create a map that the user can see (and that the other components can interact with) which specifies the exact mapping that the plugin calculated
    - Map is useful because user who starts the experiment may leave, and another user may need to stop it; good to have the state stored in the model
- When Experiment is stopped, the map is removed from the model and the results of the experiment are returned to the user and saved in the model
  - As **assets**, since these results may be large, don't want to bog down the UI by loading a lot of data that may not be used; load it on demand.



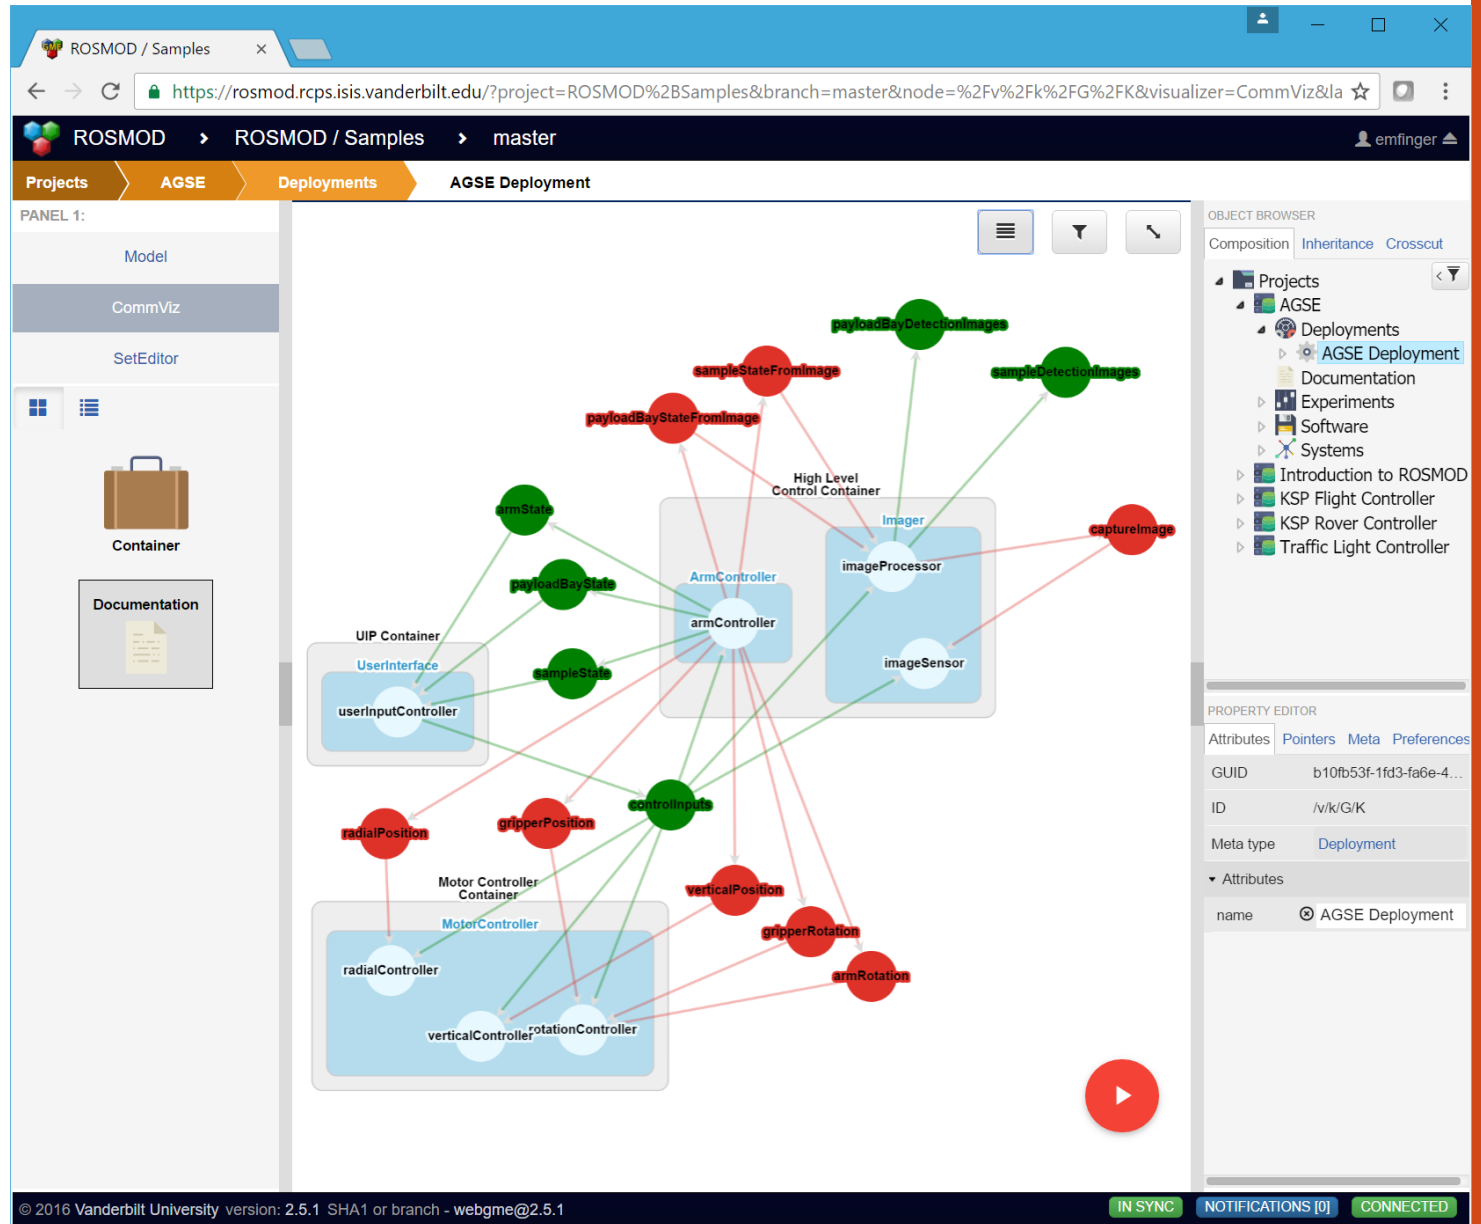
# Implementation: WebGME

- Results Visualization is important for distributed systems
  - If users have to look through tons of text logs from different processes on different nodes, they will not use your system.
  - Visualization lowers the difficulty and time it takes to find execution errors in your code/system
- Interactivity is key
  - Static plots look nice, but have limited utility when actually analyzing or debugging the system.
  - Need easy methods for users to massage the plots/data into something more meaningful for their current context.
    - Remove extra plots / data
    - Zoom x/y/x&y
    - Pan



# Implementation: WebGME Components

- Users want to know what configuration is actually running in a deployment
  - Since not all software components may be used, and the components in use may not be correctly configured
- The actual deployment may be large and difficult to visualize with many connections
  - So again, interactivity is key
- Being able to let the user select their current *context/focus* is important as the scale of the models/systems increases





# Let's look at ROSMOD and its Code!

ROSMOD: `rosmod.rcps.isis.vanderbilt.edu`

CODE: `github.com/rosmod/webgme-rosmod`

(note: some components, like the CodeEditor, are dependencies maintained in their own repositories, just look at package.json to figure out where they come from)

# Lessons Learned

- Meta specification is an evolving creature, depending on the changing needs of your platform, the target domain(s), and user experience
  - Iterative testing between over-generalization and over-dependence on implementation specifics
- As always: *documentation is important*
- Can't always just rely on WebGME built-in components
  - But also don't try to just make everything from scratch; many libraries exist that can help you do what you want to do
- The line between platform and (meta)model may shift over time and in some cases is actually more of a gray area
- When developing your code, better to err on the side of generalization / re-use since you don't know how your meta might change or what other projects may want to replicate your functionality
  - Be nice to the open-source community, since they've been so nice to you 😊
- Standardize your coding style and make use of libraries when possible to keep your code readable, e.g. **Q** for *promises* (instead of relying on callback functions)

# Thank you!

Questions?

Many acknowledgements to Brian Broll and Patrick Meijer, who were both quite patient and helpful in teaching me JS and WebGME.